

Assessment *After* AI.

The machine is in the room. The question is no longer how to keep it out, but what to measure now that it is here.

ABOUT THIS PAPER

A short argument for a different kind of assessment, written for the people who have to defend the credential.

Assessment has a problem it cannot proctor its way out of. When any person with a browser can produce a competent answer to almost any question on demand, the answer stops being evidence of anything. This paper is about what replaces it: not a better way to keep the machine out, but a better account of what a capable human looks like with the machine in the room.

It is written for academic leaders, assessment designers, credentialing bodies, and employers who suspect that both available responses, prohibition and surrender, are the same mistake wearing different clothes.

Stima Prova is the assessment instrument we are building on this argument. This paper is the argument itself.

01	The why	03
02	The false binary	04
03	The philosophy	05
04	The moat	06
05	In practice	07
06	Principles	08

SECTION 01

The exam tested a world that has *ended*.

For a century, assessment rested on one quiet assumption: the work and the worker are the same thing.

The essay, the problem set, the timed exam, the take-home project. Every format assumed that if a person produced the artifact, the person possessed the ability. The artifact was the evidence. Grade the artifact and you had graded the mind.

That assumption has broken, and it will not be repaired. A polished analysis now proves that a polished process existed somewhere. It says very little about whose judgment drove it, whose questions shaped it, or whose standards rejected the first three drafts. The artifact has been severed from the worker, and every instrument built on the artifact has been severed along with it.

Institutions feel this before they can articulate it. Grades inflate while confidence deflates. Employers discount credentials they once trusted. Students reasonably ask what, exactly, is being measured, and the honest answer is: a world in which knowledge work was done alone, with a library and a pen.

When any answer can be generated, answers stop being evidence. Judgment becomes the evidence.

The mistake would be to treat this as a temporary disruption, something to survive until detection tools catch up or norms settle. It is not a disruption. It is the new condition. Knowledge work from now on is collaboration with machines, and assessment that ignores this is not conservative. It is simply measuring the wrong thing, carefully.

SECTION 02

Ban it, or surrender to it. *Both are the same mistake.*

The institutional response has split into two camps. Both accept the artifact as the unit of evidence. That shared assumption is the error.

Camp one prohibits. Lock the browser, watch the webcam, patrol the prompt. This preserves the appearance of measurement by recreating, under fluorescent light, conditions that exist nowhere in working life. It is also an arms race with a foregone conclusion: every model generation makes generated work cheaper and detection less reliable. The proctoring camp is not defending rigor. It is defending a venue.

Camp two surrenders. If you cannot beat it, stop measuring. Grade participation, grade reflection, grade the process narrative, lower the stakes until nothing is at stake. This is kinder and more honest than camp one, and it produces credentials that certify attendance in an era. The employer reading them learns nothing, and knows it.

The question was never "how do we stop them from using it?" That question belongs to the old world. The question of the new one is: **what does a capable person look like when the machine is in the room?** Answer that, and assessment does not need to ban the machine or bow to it. It can finally do its job.

*Proctoring measures a world that no longer exists.
Surrender measures nothing at all.*

There is a third path, and it begins by changing the object of measurement. Not the artifact. The collaboration that produced it.

SECTION 03

The skill is not producing the answer.
It is *directing* the machine.

Give the same machine to a weak candidate and a strong one and you will not get the same result. That difference is the thing worth measuring, and it is visible, observable, and scorable.

Watch a skilled professional work with an AI system and a pattern emerges. They decompose before they delegate. They interrogate fluently instead of admiring it. They keep the decisions that matter and hand over the rest. And at the end, they can stand behind every line of the final work, including the lines the machine drafted, because nothing survived that they did not approve.

We score that pattern along three dimensions:

TABLE 1 · THE THREE DIMENSIONS OF POST-AI COMPETENCE

DIMENSION	THE QUESTION IT ANSWERS	WHAT IT LOOKS LIKE IN THE WORK
Direction	What did they ask, and in what order?	Task decomposition, delegation strategy, knowing what to hand the machine and what to keep. Weak candidates paste the brief. Strong candidates shape it.
Judgment	What did they accept, and what did they catch?	Verification behavior, error detection, the moment a fluent and wrong answer gets challenged. Weak candidates admire output. Strong candidates audit it.
Ownership	What did they override, and can they defend it?	Accountability under questioning. The final artifact contains nothing the candidate cannot explain, justify, or reproduce live.

None of these dimensions are new. Supervisors have read them in junior colleagues for decades: the one who asks better questions, the one who catches the confident error, the one who owns the outcome. What is new is that these are no longer soft signals around the real assessment. They are the assessment. The viva, the oldest and most AI-resistant format we have, returns in updated form: defend the work, including the machine’s part in it.

SECTION 04

A prompt can write a task. It cannot write a *measurement*.

Anyone can now ask a model to design an AI-collaborative assessment. We have. The output is plausible, fluent, and unvalidated. When generating the task becomes free, the task stops being the product.

The value of an instrument was never the worksheet. It is everything around it that makes a score mean something. Three things, plus one that compounds:

First, the construct map: knowing what to measure. An assessment is only as good as its answer to “what does mastery actually look like in this field?” Our answer comes from evidence, not imagination. A dedicated research program analyzes large corpora of real practice, job specifications, professional discourse, work artifacts, and performance language, and extracts from them the capabilities that actually distinguish strong practitioners, in the field’s own words. A model asked to invent an assessment measures what it imagines the skill to be; a corpus-derived one measures what the evidence says it is.

Second, calibrated scoring: making the grade defensible. Distinguishing good machine-direction from bad requires rubrics anchored to observed expert behavior, raters calibrated against one another, and validity treated as an argument that must be defended with evidence, not asserted.^[1] A generated rubric is fluent. A calibrated rubric is true.

Third, stability over time: a score must survive the next model. Every model release silently changes task difficulty. A task that discriminated well last year may be trivial after the next release. We version tasks, re-equate against model generations, and can say with a straight face that a 72 in 2026 means a 72 in 2028. Without equating, every score expires with the model it was taken under.

The task is now free. The measurement is the product.

And then the compounding layer. Every assessment administered adds outcome data: which collaboration behaviors predicted which real-world results. The loop sharpens the map, recalibrates the rubrics, and deepens the moat with every cohort. It cannot be prompted into existence. It can only be run.

SECTION 05

What it looks like in *practice*.

A worked example. Call it the audit brief.

The candidate receives a realistic, messy brief from the target domain: evaluate a supplier contract with three buried problems, diagnose why a process is failing from partial data, draft a policy response under competing constraints. An AI assistant is available throughout, the same class of tool the candidate will use in the job. Ninety minutes of work. Fifteen minutes of defense.

Three evidence streams are scored together:

TABLE 2 · SCORING SIGNALS IN A POST-AI ASSESSMENT

SIGNAL	WHAT WE WATCH	WEAK SIGNAL	STRONG SIGNAL
The transcript	The full record of the collaboration: prompts, revisions, challenges	Pastes the brief, accepts the first fluent answer, no verification moves	Decomposes the problem, sequences requests, probes and challenges output
The artifact	The final work product, graded against domain standards	Polished surface, buried errors, no evidence of prioritization	Correct where it matters, trade-offs made explicit and defensible
The defense	Fifteen minutes of live questioning on any choice, including the machine's	Cannot explain why an answer was accepted; ownership dissolves under questions	Explains every kept and overridden choice; could do it again, differently

Note what is absent: no detection software, no locked browser, no attempt to prove the machine was unused. The machine is used, openly, by every candidate. The assessment does not ask whether you can beat the tool. It asks whether you can lead it. And the format extends outward, to admissions, hiring, credential renewal, and reskilling. Anywhere the question is "what can this person actually do," the honest modern version is "what can this person do **with**."

SECTION 06

Principles for the *post-AI* era.

01 The machine is in the room.

Design for that room. Every format that pretends otherwise measures a fiction.

02 Never grade the artifact alone.

Grade the collaboration that produced it: the direction, the judgment, the ownership.

03 Direction is a skill. Measure it.

What a candidate asks, in what order, tells you more than what they receive.

04 Fluency is not competence.

The most dangerous output is confident and wrong. Score the verification, not the admiration.

05 Ownership is visible under questioning.

Ask. Fifteen live minutes reveal more than fifteen polished pages.

06 A score must survive the next model release.

Version tasks. Equate difficulty. A measurement that expires with the model was never a measurement.

07 Assessment should raise esteem, not just rank it.

Stima means both. An instrument that only sorts people has failed half its job.

08 Prohibition is nostalgia. Measurement is respect.

Treat candidates as the professionals you say you are preparing them to become.

FURTHER READING

Messick, S. (1989). Validity. In R. L. Linn (Ed.), *Educational measurement* (3rd ed., pp. 13–103). American Council on Education & Macmillan.

Wiggins, G. (1990). The case for authentic assessment. *Practical Assessment, Research, and Evaluation*, 2(2).

Bloom, B. S. (1984). The 2 sigma problem: The search for methods of group instruction as effective as one-to-one tutoring. *Educational Researcher*, 13(6), 4–16.

WHAT WE ARE BUILDING

Credentials that survive contact with the *real world*.

Stima Prova is the assessment instrument built on this argument: evidence-derived construct maps, calibrated scoring, versioned against every model release, sharpened by every cohort. We are onboarding early institutional partners through 2026.

If your institution needs its credentials to mean something in the era of human and machine collaboration, we should talk before your next intake does.

STIMA PROVA • PRE-LAUNCH

EARLY ACCESS AND PARTNERSHIPS: [MYSTIMA.IO](https://mystima.io)

POSITION PAPER NO. 01 • © 2026 STIMA

Know your worth. Prove it.

stima •